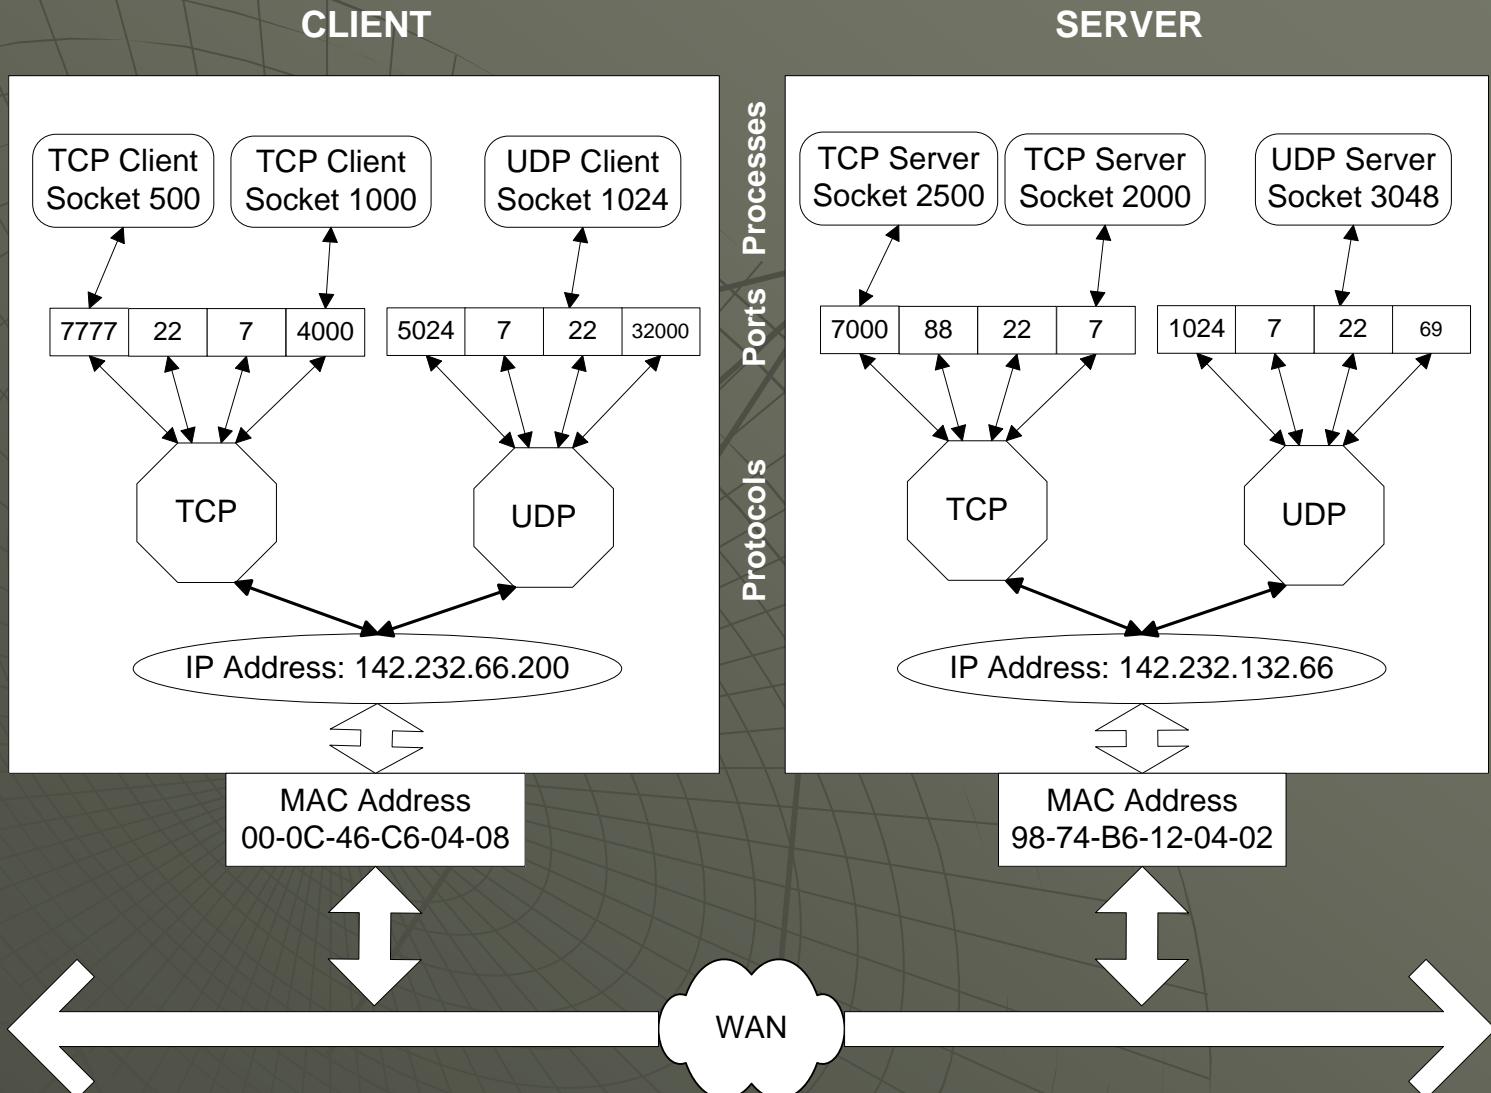


TCP/IP Programming Interface

- **Network programming involves the implementation of “server” and “client” applications.**
- **We will use the Berkeley socket API to design “server” and “client” applications**

TCP/IP Model



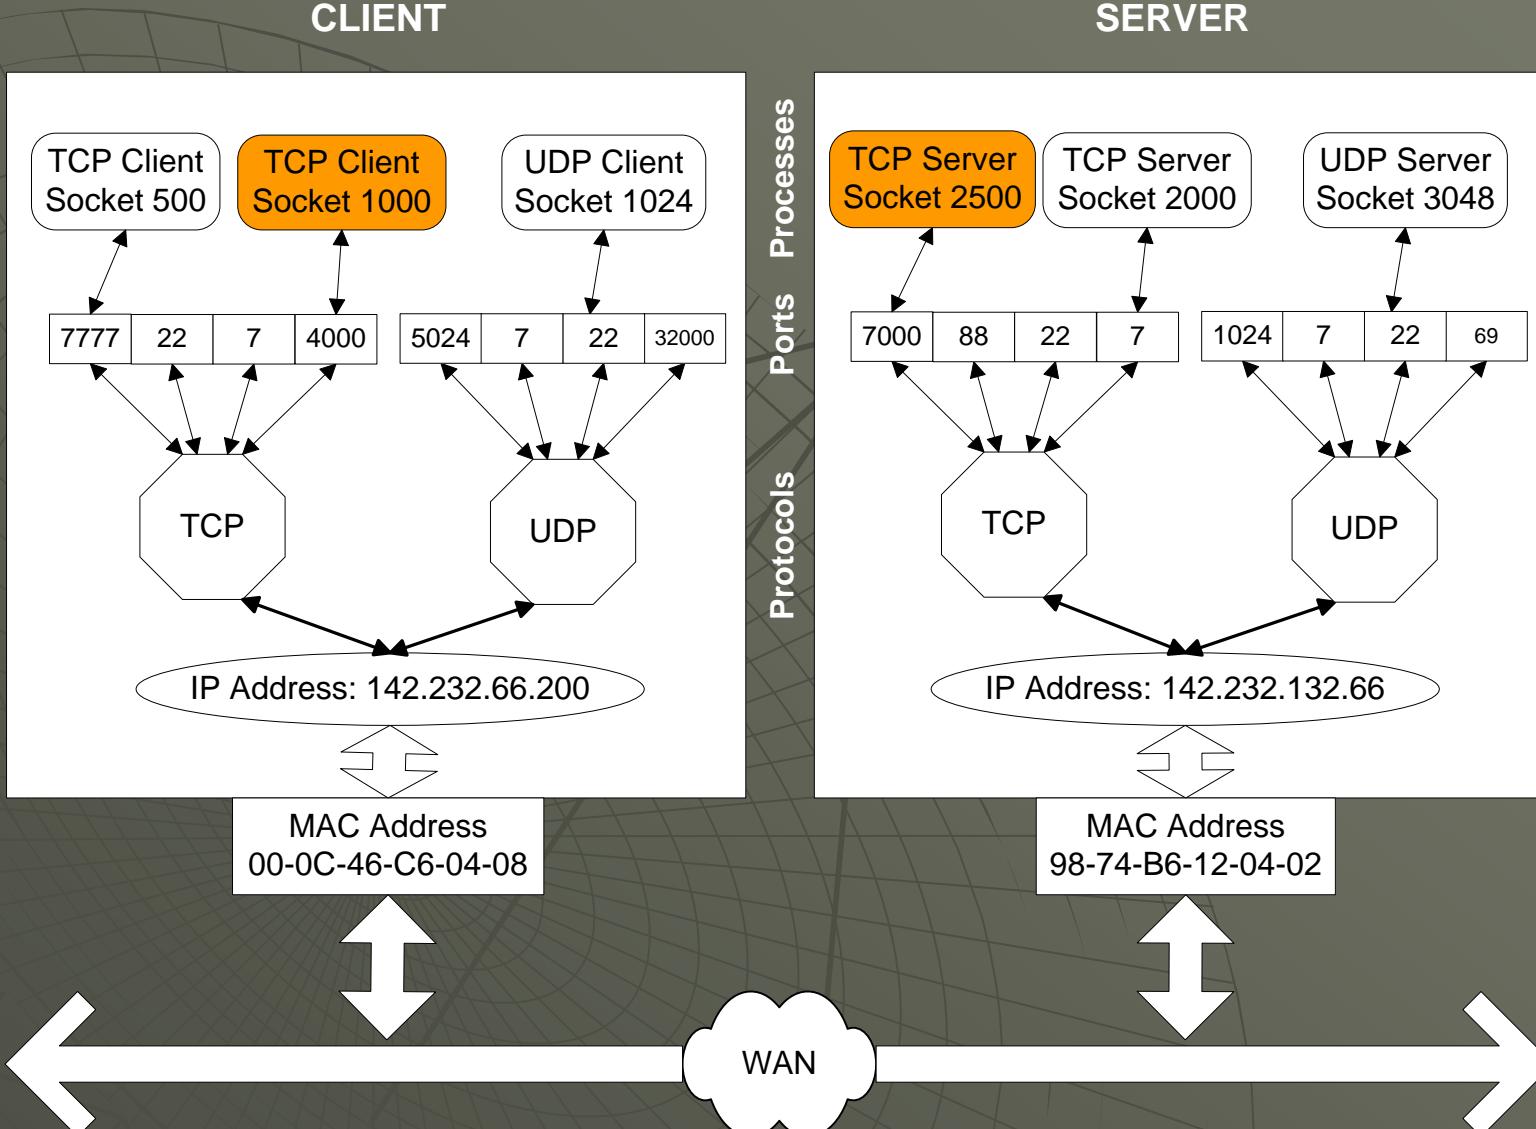
Socket

- ◆ Generalization of the **UNIX file access system** designed to incorporate **network protocols**.
- ◆ **Communications channel** between a pair of '**sockets**'
- ◆ A socket is defined by a group of four integers:
 - The remote host identification number or address
 - The remote host port number
 - The local host identification number or address
 - The local host port number

Create Socket

`socket(family, type, protocol)`

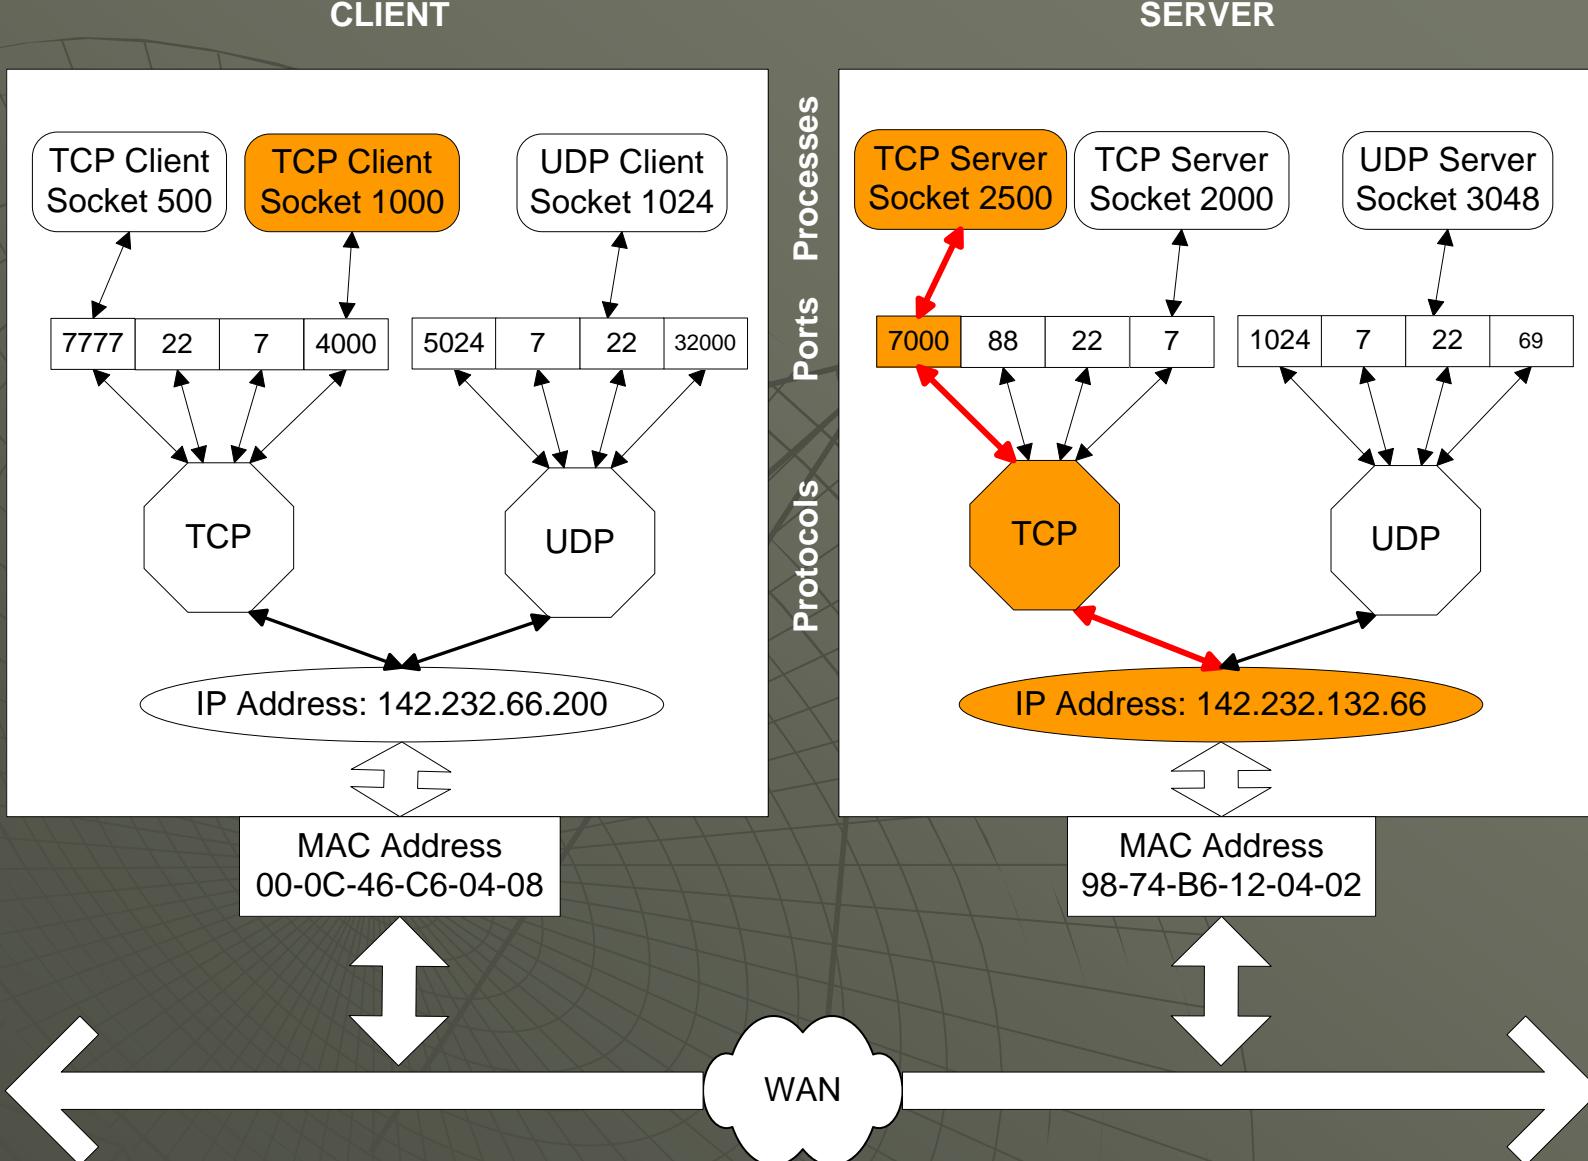
- ◆ family: the protocol family
(e.g. **PF_INET**, **PF_UNIX**)
- ◆ type: the abstract type of the communication desired
(e.g. **SOCK_STREAM**, **SOCK_DGRAM**)
- ◆ protocol: specific protocol desired
(e.g. **TCP** or **UDP**)
- ◆ Returns an socket file descriptor (int)



Bind Socket with Address - server

**bind (socket,
sockaddr,
sockaddrlen)**

- ◆ socket: socket descriptor
- ◆ sockaddr: **pointer** to the **address** to which the socket should be bound.
- ◆ sockaddrlen: the **size** of the **address**.
- ◆ Returns -1 and sets **errno** for error



listen() - server

Marks socket as 'listening' and sets the maximum number of listen queue

int listen(int s, int backlog)

- ◆ s: socket
- ◆ backlog: max length of the queue of unprocessed connection requests

Addressing

- ◆ **common framework** for all addresses to support multiple protocol families
- ◆ In the **Internet family**, transport addresses are **6 octets** long:
 - **4 octets** for the **Internet address**
 - **2 octets** for the **port number**.
- ◆ **<netinet/in.h>** defines the appropriate structures to use for the Internet family

```
struct sockaddr_in
{
    short sin_family;
    unsigned short sin_port;
    struct in_addr sin_addr;
    unsigned char sin_zero[8]; /* unused */
}
```

sin_family: address family

sin_port: 16-bit port number

sin_addr: 32-bit Internet address

```
struct in_addr
{
    u_long s_addr;           /* 32-bit IP Address */
}
```

Get Server Information

`gethostname(host_name)`

- ◆ Maps the request into a **DNS (Domain Name System)** query that it sends to the **resolver** running on the local machine.
- ◆ Returns a **pointer** to a **hostent structure** that contains the requested addresses.
- ◆ The related function **gethostbyaddr()** takes an **Internet address** as an argument instead.

```
struct hostent
{
    char *h_name;
    char **h_alias;
    int h_addrtype;
    int h_length;
    char **h_addr_list;
};
```

h_name: the official, fully qualified **domain name**

h_alias: a null-terminated list of alternate names
(aliases)

h_addrtype: **protocol family** the address belongs to

h_length: the **length** of the address

h_addr_list: a null-terminated **array of addresses**

Address Manipulation Functions

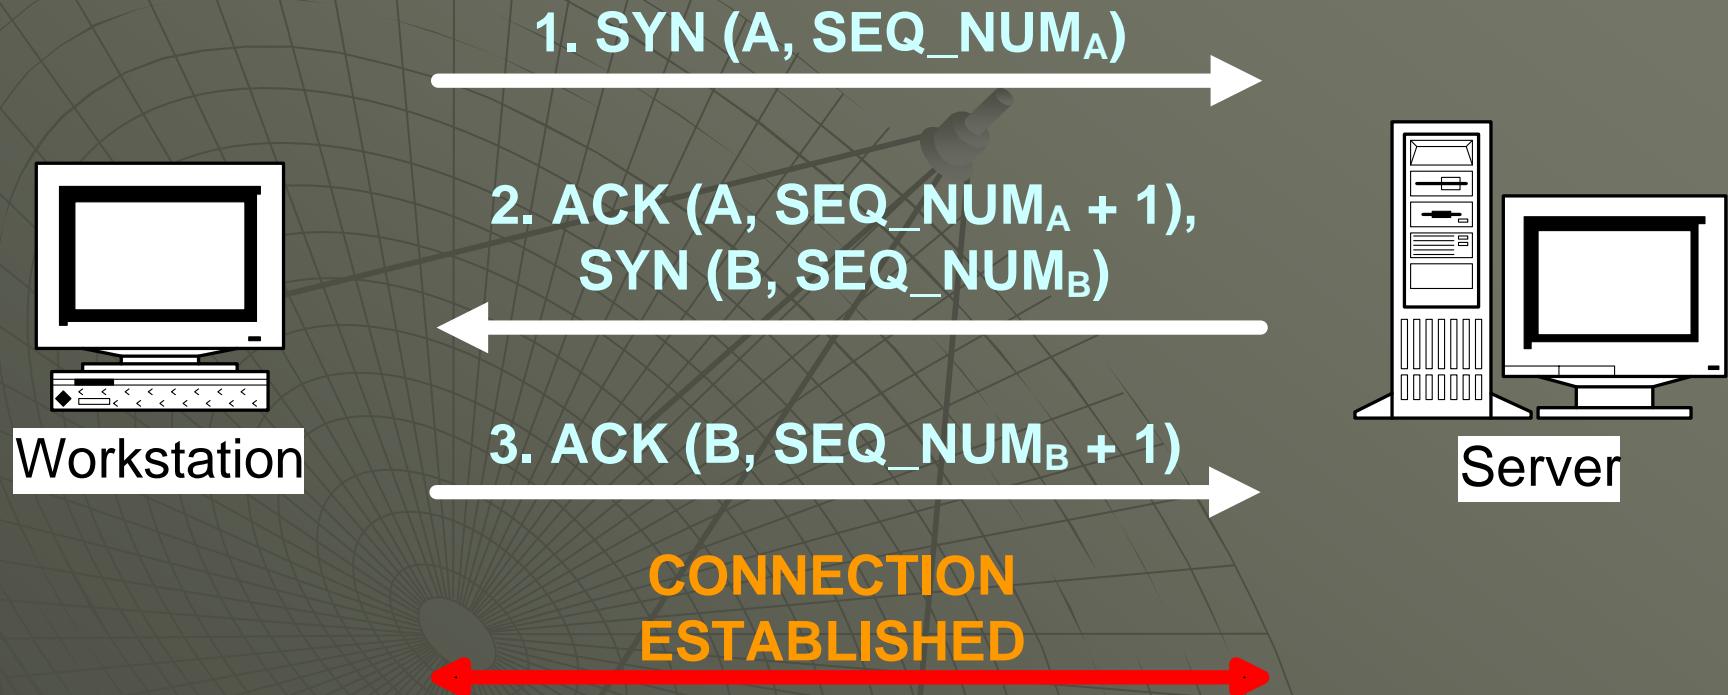
`inet_aton()` / `inet_network()`

- ◆ Takes strings representing Internet addresses in **dotted notation**
- ◆ Returns numbers suitable for use in **sockaddr_in** structures.

`inet_ntoa()`

- ◆ Takes a 32-bit IP address in **network byte order** Internet address
- ◆ Returns the corresponding address in dotted-decimal notation.

Make a Connection to Server

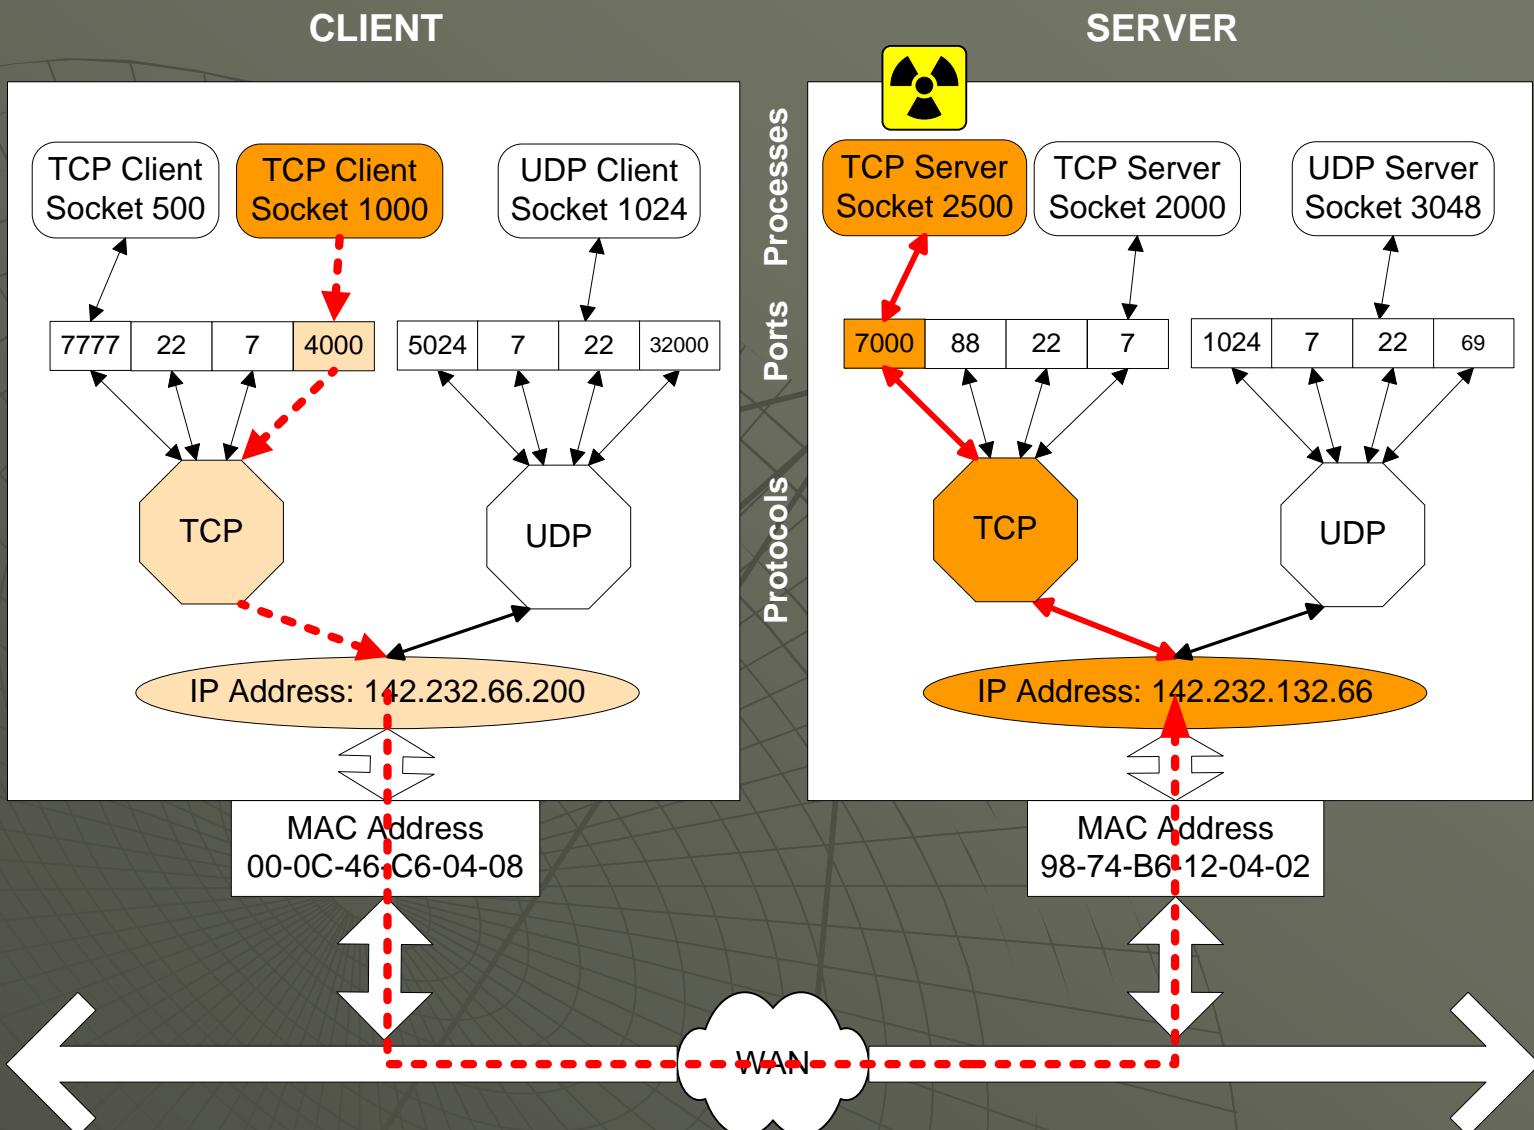


TCP 3-Way Handshake

Send Connect Request (SYN)

```
connect( int s,  
        struct sockaddr *name,  
        int namelen )
```

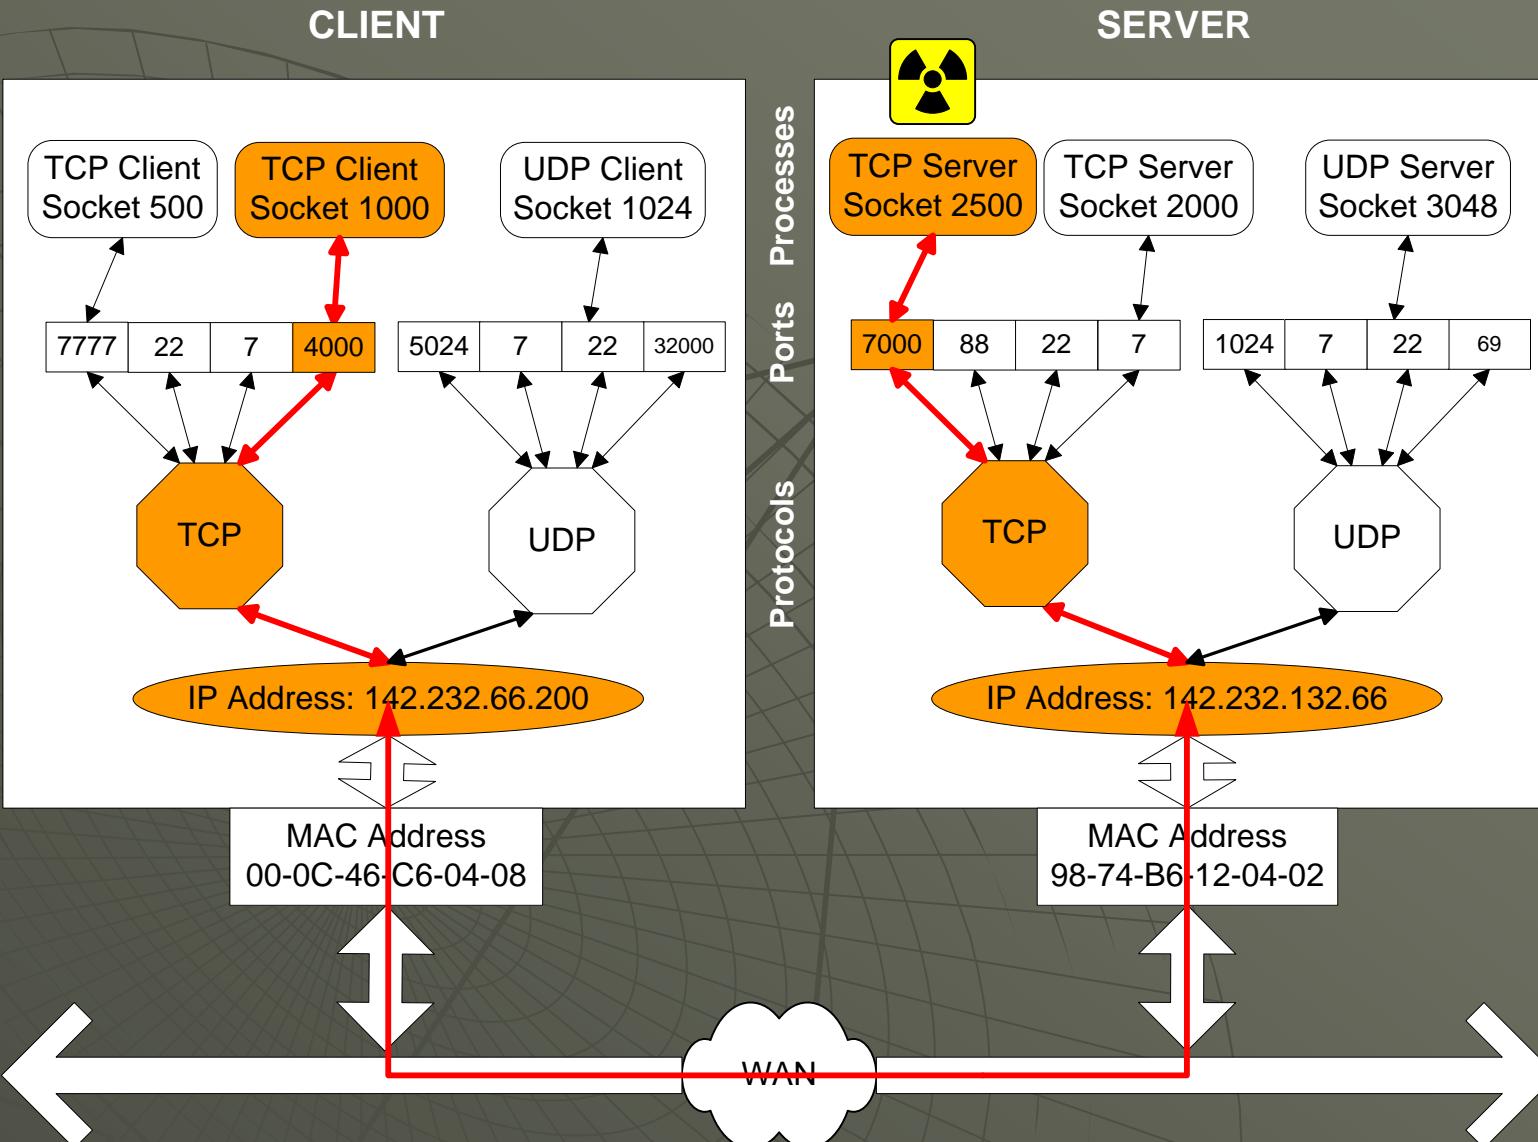
- ◆ Establishes communication with a remote entity
 - initiates a **three-way handshake** for TCP connection
 - normally only used for **SOCK_STREAM** sockets.
- ◆ s: client socket
- ◆ name: points to an area of memory containing the address information for the remote entity
- ◆ namelen: gives the length of 'name'



Accept Connect Request (SYN/ACK)

```
accept ( int s,  
        struct sockaddr *addr,  
        int *addrlen)
```

- ◆ s: base socket descriptor
- ◆ addr: points to a struct sockaddr of the remote (calling) system.
- ◆ addrlen: size of addr
- ◆ Returns a new socket descriptor that will be used for all subsequent communication with the remote host.



Send/Receive Data

send(socket, message, length, flags)
recv(socket, message, length, flags)

- ◆ message: a pointer to the data to be written.
- ◆ length: data length.
- ◆ flags: to send **flags** to the **underlying protocol**.
 - flags may be formed by ORing MSG_OOB and MSG_DONTROUTE.
- ◆ `send()` and `recv()` may only be used with **SOCK_STREAM** type sockets.

Send/Receive Data - UDP

**sendto(socket, message, length,
flags, dest, destlength)**

**recvfrom(socket, message,
length, flags, from, fromlength)**

- ◆ The two additional arguments are **pointers** to a **sockaddr** structure and its **length**.

Close Connection

shutdown (socket, how)

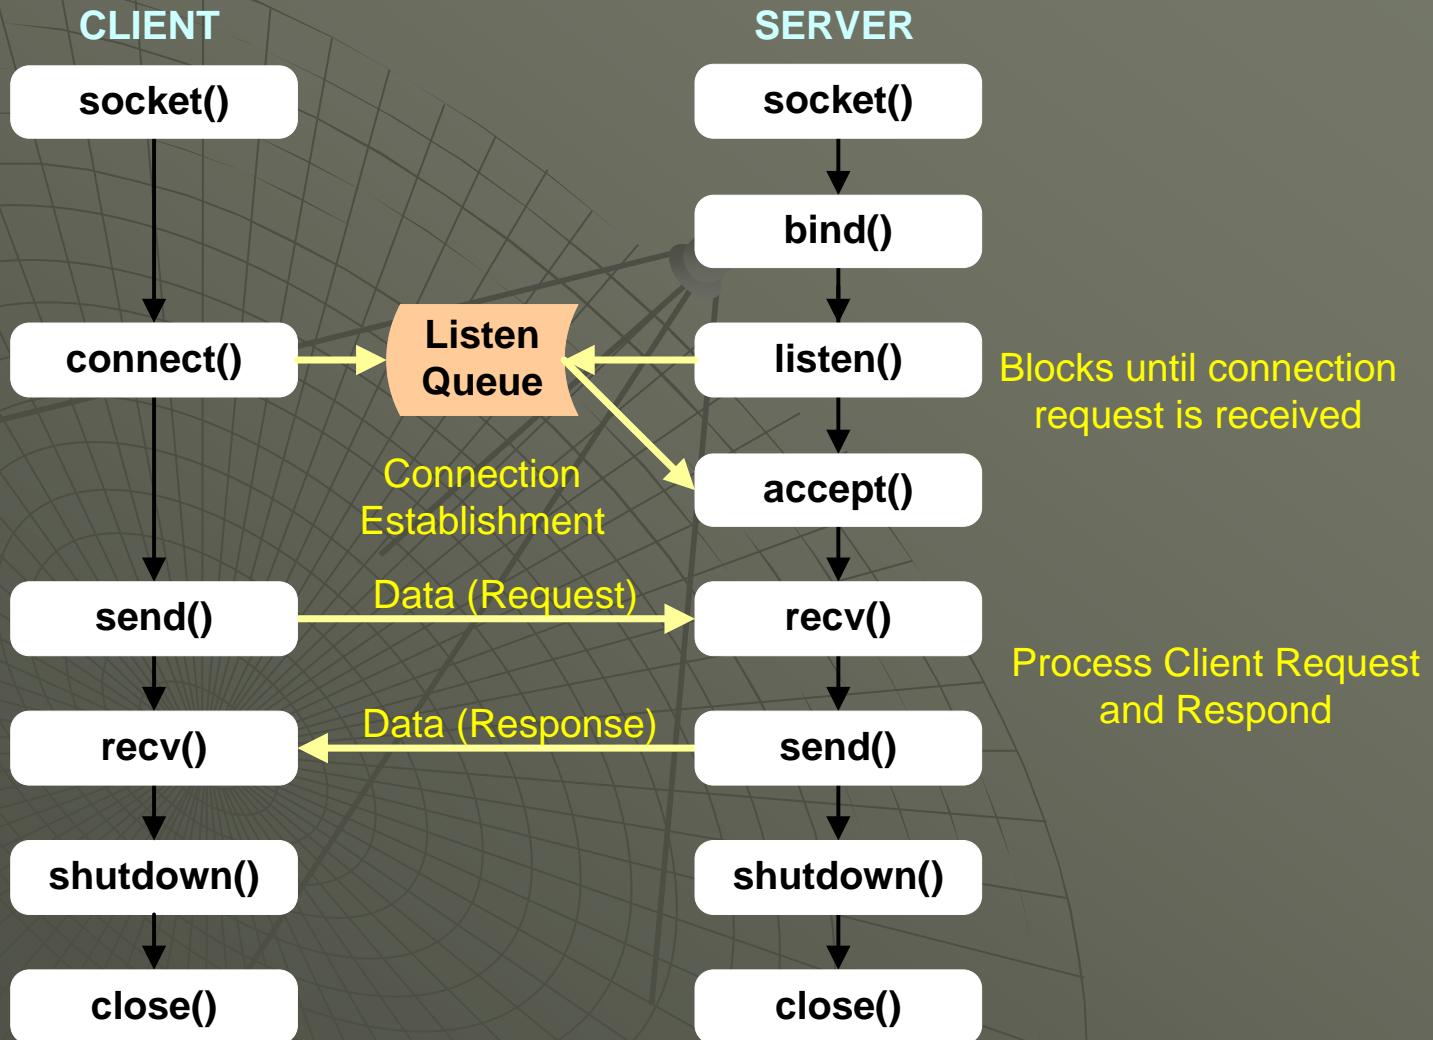
to shut down all or part of a full-duplex connection

- ◆ how: **direction** of the connection to be shutdown:
 - 0 - additional receives are shutdown.
 - 1 - additional sends are disallowed.
 - 2 - additional sends and receives will be disabled.

close(socket)

to close a socket

TCP/IP Client Server Model



UDP Client Server Model

